

Reproducible, flexible and high throughput data extraction from primary literature for meta-analyses: The metaDigitise R package

Joel L. Pick, Shinichi Nakagawa & Daniel W.A. Noble

November 3, 2017

Prepared for: Journal of Statistical Software

Abstract

Keywords: meta-analysis, data extraction, R, reproducibility, figures, images, summary statistics

Introduction

Meta-analysis is becoming increasingly common in many research fields. They make use of various summary statistics from studies to generate effect size estimates and their sampling error variance to understand whether particular experimental manipulations, or relationships between variables, show strong overall effects and attempt to explain variation among these effects (Nakagawa et al. 2017). Meta-analysis relies foremost on data extracted from primary literature – more specifically statistics that have been reported in the text or tables of research papers. However, necessary summary statistics are often presented in figures, and so, these need to be manually extracted using digitising programs to be used in the final analysis. Although there are several existing tools to perform tasks like this, these tools are not made for this specific purpose (i.e. a meta-analysis). Specifically, they do not differentiate between common plot types that are used to present data, meaning that they require a large amount of downstream data manipulation. For example, extracting often only provides the raw data or values from plots. This means that these data need to be imported into other programs to calculate correlations (if scatter plots) or errors on means (from mean error plots) need to be back calculated by hand to convert them to a common type (e.g., standard errors to standard deviations). Digitising programs often do not allow for different grouping of points to be made, and frequently do not integrate metadata (such as variable names, group names and sample sizes) at the time of data extraction. Additionally, existing stand-alone software does not allow easy import into commonly used statistical software (such as R), or try to optimize the research pipelines to facilitate editing and reproducibility. These are major issues because extracting from figures can be an incredibly time-consuming process. Furthermore, meta-analyses can have a major impact on research fields, and should be expanded upon as more research comes to light. Having a tool that facilitates reproducibility in meta-analyses will increase transparency and go a long way to resolving the reproducibility crises we are seeing in many fields.

Here, we present an interactive R package, ‘metaDigitise’, which is designed for large scale data extraction from figures, specifically catering to the needs of meta-analysts. To this end, we provide tools specific to data extraction from common plot types (mean and error plots, box plots, scatter plots and histograms, see Figure 1). ‘metaDigitise’ operates within the R environment making data extraction, analysis and export more streamlined. It also provides users with options to conduct the necessary calculations on raw data immediately after extraction so that comparable summary statistics can be obtained quickly. Summaries will condense multiple figures into data frames or lists (depending on the type of figures) and these objects can easily be exported from R, or if using the raw data, analysed in any way the user desires. Conveniently, when needing to process many figures at different times ‘metaDigitise’ will only import figures not already completed within a directory. This makes it easy

to add new figures at anytime. ‘metaDigitise’ has also been built for reproducibility in mind. It has functions that allow users to redraw their digitisations on figures, correct anything and access the raw calibration data which is written automatically for each figure that is digitised into a special folder within the directory. This makes sharing figure digitisation and reproducing the work of others simple and easy, and allows meta-analysts to update meta-analyses more easily.

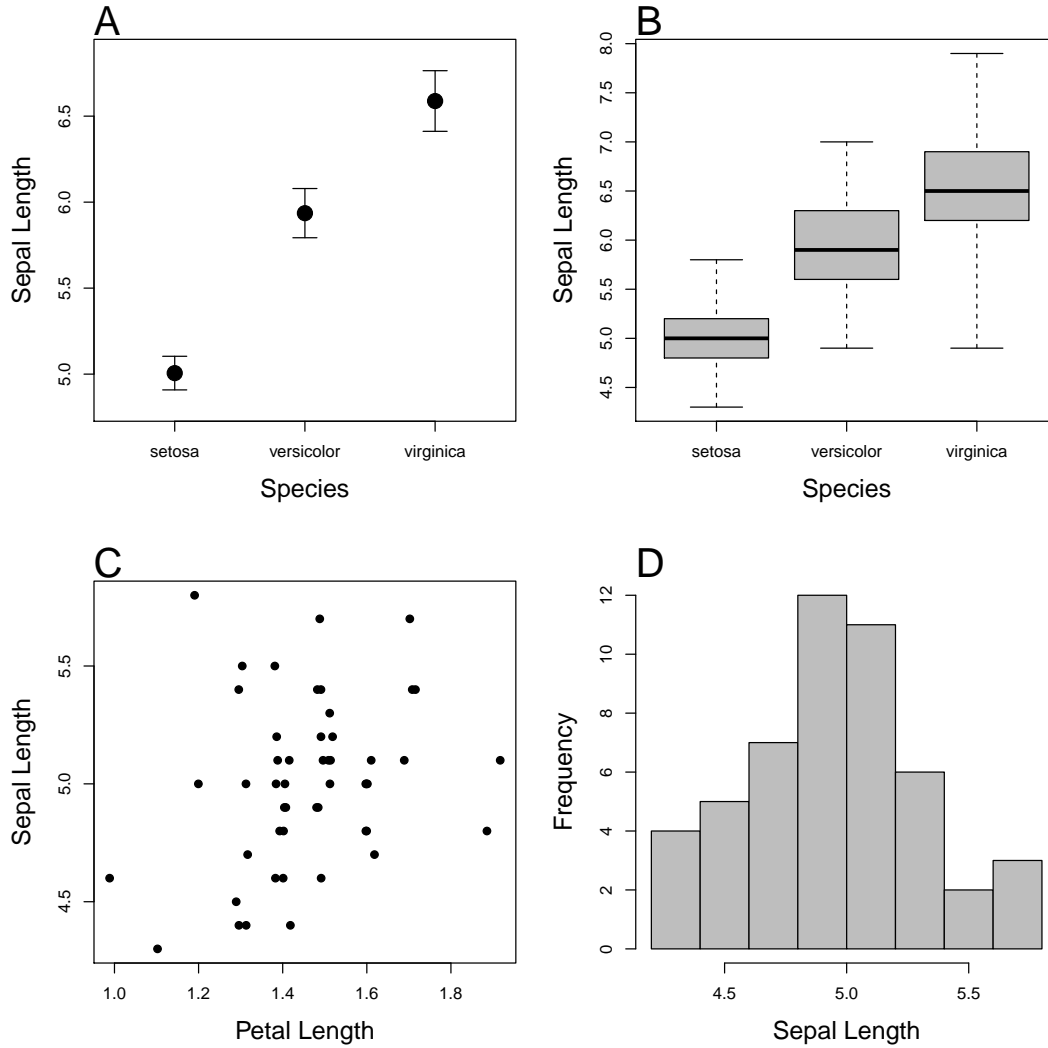


Figure 1: Four plot types that metaDigitise is designed to extract data from: A) mean and error plot, B) box plot, C) scatter plot and D) histogram. Data is taken from the iris dataset in R. A and B are plotted with the whole dataset, C and D are just the data for the species setosa.

1 Functionality

There is one main function in the package metaDigitise, which interactively takes the user through the process of extracting figures from graphs.

1.1 Bulk Digitisation

metaDigitise was created with the idea that the user would have multiple images to extract from. It therefore operates in the same way whether the user has one or multiple images. It assumes that the user has put all the figures needed for extraction in one parent directory. It first creates a new directory ‘caldat’ inside this parent directory, in which calibration files are saved (see below).

It then prompts the user to select whether they want to extract from new images, import previously extracted data or edit previously extracted data. Extracting from new images takes the user through the steps outlined below, figure by figure, giving the user the option to quit the process after each figure. If this process has already been started when the user starts metaDigitise, then it starts where the user left off. Similarly if the user adds new images to the parent directory after all other images are extracted (or part way through), the new images will be integrated into the extraction procedure. After extraction (or upon quitting the function), the extracted data will be returned.

If the process of data extraction has been finished, or the user wants to reimport summary or processed data ...

```
install.packages("devtools")
devtools::install_github("daniellnoble/metaDigitise")
library(metaDigitise)
```

1.2 Figure Rotation

Figure may have been extracted from old publications, for example from scanned images, and so are not perfectly orientated on the image. This will make the calibration of the points in the figure from the image problematic. metaDigitise allows users to rotate the image. By clicking two points on the x-axis, metaDigitise calculates the angle needed to rotate the image so the x-axis is horizontal, and rotates it. (Figure 2A,B)

Furthermore, in some figures mean and error, boxplots or histograms may be presented with horizontal bars. metaDigitise assumes that the bars are vertical, but allows the user to flip the image so that the bars are vertical (Figure 2C,D).

1.3 Calibration

metaDigitise requires the user to calibrate the axes in the figure. To do this the user is required to click on two known points on the axis in question, and then enter the value of those points in the figure (Figure 3). Using this information, metaDigitise then calculates the value of any clicked points in terms of the




Figure 2: Figure rotation. A) and B) show how non-aligned imaged can be realigned through user defined rotation. C) and D) show how figures can be re-orientated so as to aid data input.

figure axes. In the case of mean and error plots and box plots, it calibrates only the y axis (assuming the x axis is redundant). For scatter plots and histograms both axes are calibrated.

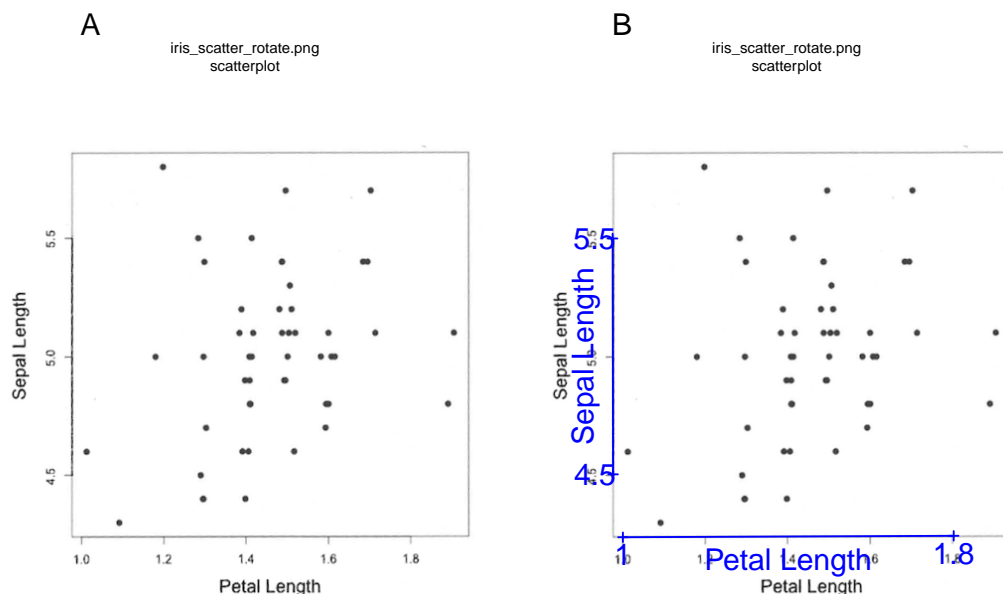


Figure 3: Axis calibration. The user defines two points on each axis and labels them according to the values shown in the figure.

1.4 Plot Types

metaDigitise recognises 4 main types of plot; Mean and error plots, box plots, scatter plots and histograms, shown in Figure 1.

1.4.1 Mean and error plots

metaDigitise prompts the user to enter groups names and allows the user to enter sample sizes, which are used in downstream processing. The user is then prompted to click on an error bar followed by the mean. Error bars above or below the mean can be clicked - sometimes one is clearer than the other. metaDigitise assumes that the error bars are symmetrical. Where the user has clicked the error is displayed in a different colour to the mean (Figure 4A). The user can subsequently add more groups or remove groups.

1.4.2 Box plots

metaDigitise prompts the user to enter groups names and allows the user to enter sample sizes, which are used in downstream processing. The user is then prompted to click on the maximum, upper quartile, median, lower quartile and minimum. metaDigitise will check that the maximum is greater than the minimum, and return an warning if that is not the case. The user can subsequently add more groups or

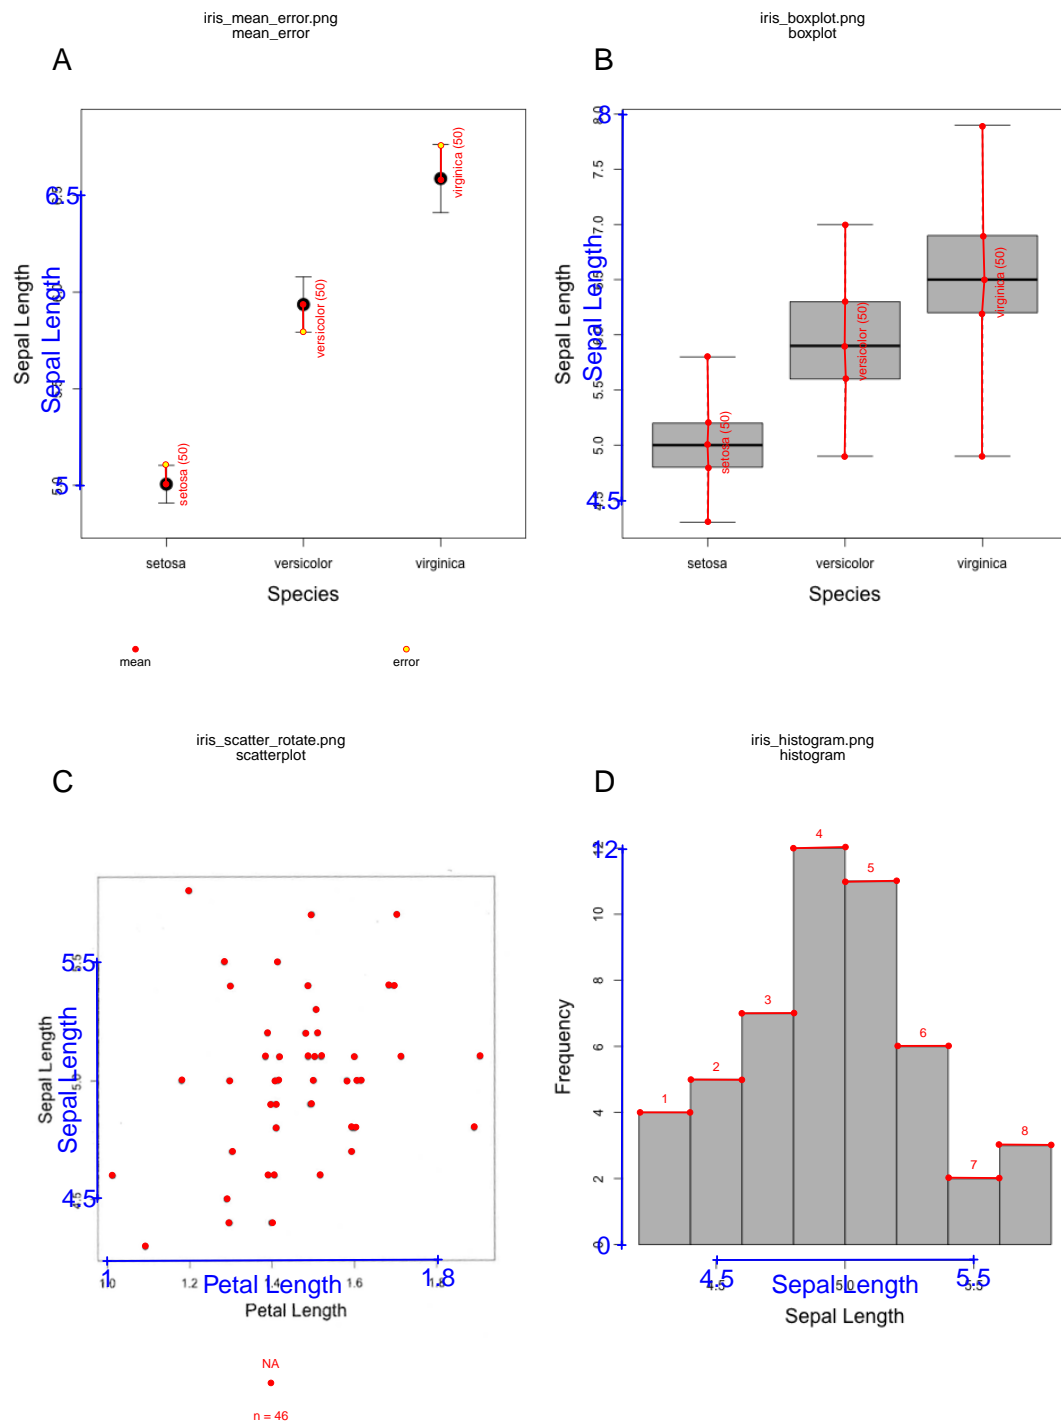


Figure 4: Demonstration of data extraction from different plot types

remove groups.

1.4.3 Scatter plots

metaDigitise prompts the user to enter groups names and then to click on points. Points added by mistake can be deleted. The user can subsequently add groups, edit groups (add or remove points) or delete groups. Different groups are plotted in different colours and shapes, with a legend at the bottom of the figure (Figure 4C). The x and y coordinates are returned.

1.4.4 Histograms

metaDigitise prompts the user to enter groups names and then to click on the top corners of each bar. bars can subsequently be deleted. MetaDigitise then calculates the midpoint and the frequency for each bar.

1.5 Summary Data

From all plot types, metaDigitise summarises the data from the figure to a mean, standard deviation and sample size, for each identified group. These are the summary statistics needed to create all relevant statistics for a meta-analysis. In the case of scatter plots, metaDigitise also returns the correlation coefficient between the points within each identified group.

2 Directory Structure and Reproducibility

The ‘metaDigitise’ package is quite flexible. Users can extract single figures (if this is all they have) using the ‘metaDigitise’ function with a path name to the directory with the file. However, often many figures need extracting from a single paper or set of papers. ‘metaDigitise’ will also handle these situations seamlessly by simply cycling through all figures within a directory. This is useful because it expedites digitising figures as it prevents users from having to constantly specify the directories and / or paths where files are stored. ‘metaDigitise’ essentially will bring up each figure within a folder automatically and allow the user to click and enter the relevant information about a figure as they go. This information is then all stored in a data frame or list at the end of the process, saving quite a bit of time. Users can stop mid-way through a folder by simply exiting after the last plot they have digitised. The data from completed figures will automatically be written to the ‘caldat/’ folder for later use and editing, should the user need to do this.

‘metaDigitise’ can work on a directory with figures (currently .png, .jpg, .tiff, .pdf images can be used) from many different papers and that are of different types. However, users can get creative in how they set up the directories of figures to facilitate extraction. For example, one might have 3–4 figures

from a single paper that need extracting and the user may want to focus on a single paper at a time while the information about a paper is on hand. This could be done by simply setting up a file structure as follows and then using ‘metaDigitise’ with path names (i.e., directory) for each papers folder:

```
* Main project directory
  + Paper1_P1
    + Figure1.png
    + Figure2.png
    + Figure3.png
  + Paper2_P2
    + Figure1.png
    + Figure2.png
    + Figure3.png
```

An alternative directory structure (and probably the most flexible) would be to simply have a set of different figures with an informative and relevant naming scheme to make it easy to identify the paper and figure the data come from. This cuts out the need to change directories constantly. For example the directory structure could look like:

```
* Main project directory
  + FiguresToExtract
    + P1_Figure1_trait1.png
    + P1_Figure2_trait2.png
    + P1_Figure3_trait3.png
    + P2_Figure1_trait1.png
    + P2_Figure2_trait2.png
    + P2_Figure3_trait3.png
```

The above directory structure is probably the easiest in combination with a clear and unambiguous naming scheme for each figure. Even if only figures from a single paper are digitised, one paper at a time, an overall figure directory will work perfectly because ‘metaDigitise’ will only cycle through incomplete figures, so figures can be added at anytime.

Nonetheless, how users set up their directory is really up to them. However, it is important for users to think carefully about reproducibility at this stage. Would they like to share the entire project folder with colleagues? or would they prefer to simply share the image folder? The answers to these questions are important because relative path names are stored in ‘metaDigitise’ objects, meaning that the directory structure (currently at least) needs to be the same for colleagues to re-load previously

digitised objects. Therefore, the working directory needs to be set to either the "main project directory" or the "FiguresToExtract" directory and relative path names (e.g., `dir = "./FiguresToExtract/"` if setting the working directory to "Main project directory") used as opposed to absolute paths. This facilitates reproducibility should colleagues want to see the digitisations or integrate previously digitised figures with newly updated figures. All users need to provide them with is the project or image folder containing the `caldat/` folder and the images.

2.1 Plotting data extraction

2.2 Editing data extraction

3 Example

4 Testing